

Project: Predict Sleep Health and Lifestyle

Feature Description

- Person ID: An identifier for each individual.
- Gender: The gender of the person (Male/Female).
- Age: The age of the person in years.
- Occupation: The occupation or profession of the person.
- Sleep Duration (hours): The number of hours the person sleeps per day.
- Quality of Sleep (scale: 1-10): A subjective rating of the quality of sleep, ranging from 1 to 10.
- Physical Activity Level (minutes/day): The number of minutes the person engages in physical activity daily.
- Stress Level (scale: 1-10): A subjective rating of the stress level experienced by the person, ranging from 1 to 10.
- BMI Category: The BMI category of the person (e.g., Underweight, Normal, Overweight).
- Blood Pressure (systolic/diastolic): The blood pressure measurement of the person, indicated as systolic pressure over diastolic pressure.
- Heart Rate (bpm): The resting heart rate of the person in beats per minute.
- Daily Steps: The number of steps the person takes per day.
- Sleep Disorder: The presence or absence of a sleep disorder in the person (None, Insomnia, Sleep Apnea).

Details about Sleep Disorder Column:

- None: The individual does not exhibit any specific sleep disorder.
- Insomnia: The individual experiences difficulty falling asleep or staying asleep, leading to inadequate or poor-quality sleep.
- Sleep Apnea: The individual suffers from pauses in breathing during sleep, resulting in disrupted sleep patterns and potential health risks.

Import Library

```
In [1]: import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from statistics import mean
sns.set_theme(color_codes=True)

from sklearn.model_selection import train_test_split, KFold
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import HistGradientBoostingClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score
from imblearn.over_sampling import SMOTE

from imblearn.over_sampling import RandomOverSampler

import warnings
warnings.filterwarnings("ignore")
```

Function

```
In [3]: def sexo(texto):
        if texto == 'Male':
            return(0)
        else:
            return(1)

def BMI_Cat(texto):
    if texto == 'Normal':
        return(0)
    elif texto == 'Overweight':
        return(1)
    elif texto == 'Normal Weight':
        return(2)
    else:
        return(3)

In [4]: def target(texto):
        if texto == 'None':
            return(0)
        elif texto == 'Sleep Apnea':
            return(1)
        else:
            return(2)
```

Import Dataset

```
In [5]: df_health = pd.read_csv('Sleep_health_and_lifestyle_dataset.csv')
```

```
In [6]: df_health.head()
```

Out[6]:

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Blood Pressure	Heart Rate	Daily Steps	Sleep Disorder	
	0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	126/83	77	4200	None
	1	2	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75	10000	None
	2	3	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75	10000	None
	3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85	3000	Sleep Apnea
	4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85	3000	Sleep Apnea

Data Preprocessing Part 1

```
In [7]: # Tamanho do dataset
df_health.shape
```

```
Out[7]: (374, 13)
```

```
In [8]: # Conferindo os tipos dos dados
df_health.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 374 entries, 0 to 373
Data columns (total 13 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Person ID              374 non-null   int64
1   Gender                 374 non-null   object
2   Age                   374 non-null   int64
3   Occupation             374 non-null   object
4   Sleep Duration         374 non-null   float64
5   Quality of Sleep       374 non-null   int64
6   Physical Activity Level 374 non-null   int64
7   Stress Level           374 non-null   int64
8   BMI Category           374 non-null   object
9   Blood Pressure         374 non-null   object
10  Heart Rate              374 non-null   int64
11  Daily Steps            374 non-null   int64
12  Sleep Disorder         374 non-null   object
dtypes: float64(1), int64(7), object(5)
memory usage: 38.1+ KB
```

```
In [9]: # Validando a quantidade de valores unicos nas colunas categoricas
```

```
df_health.select_dtypes(include='object').nunique()
```

```
Out[9]: Gender      2
Occupation    11
BMI Category   4
Blood Pressure 25
Sleep Disorder 3
dtype: int64
```

```
In [10]: # Observando os valores categoricos
df_health['Occupation'].unique()
```

```
Out[10]: array(['Software Engineer', 'Doctor', 'Sales Representative', 'Teacher',
                'Nurse', 'Engineer', 'Accountant', 'Scientist', 'Lawyer',
                'Salesperson', 'Manager'], dtype=object)
```

```
In [11]: df_health['Blood Pressure'].unique()
```

```
Out[11]: array(['126/83', '125/80', '140/90', '120/80', '132/87', '130/86',
                '117/76', '118/76', '128/85', '131/86', '128/84', '115/75',
                '135/88', '129/84', '130/85', '115/78', '119/77', '121/79',
                '125/82', '135/90', '122/80', '142/92', '140/95', '139/91',
                '118/75'], dtype=object)
```

Exploratory Data Analysis (EDA)

```
In [12]: # Resumo estatístico
df_health.describe()
```

	Person ID	Age	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	Heart Rate	Daily Steps
count	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000	374.000000
mean	187.500000	42.184492	7.132086	7.312834	59.171123	5.385027	70.165775	6816.844920
std	108.108742	8.673133	0.795657	1.196956	20.830804	1.774526	4.135676	1617.915679
min	1.000000	27.000000	5.800000	4.000000	30.000000	3.000000	65.000000	3000.000000
25%	94.250000	35.250000	6.400000	6.000000	45.000000	4.000000	68.000000	5600.000000
50%	187.500000	43.000000	7.200000	7.000000	60.000000	5.000000	70.000000	7000.000000
75%	280.750000	50.000000	7.800000	8.000000	75.000000	7.000000	72.000000	8000.000000
max	374.000000	59.000000	8.500000	9.000000	90.000000	8.000000	86.000000	10000.000000

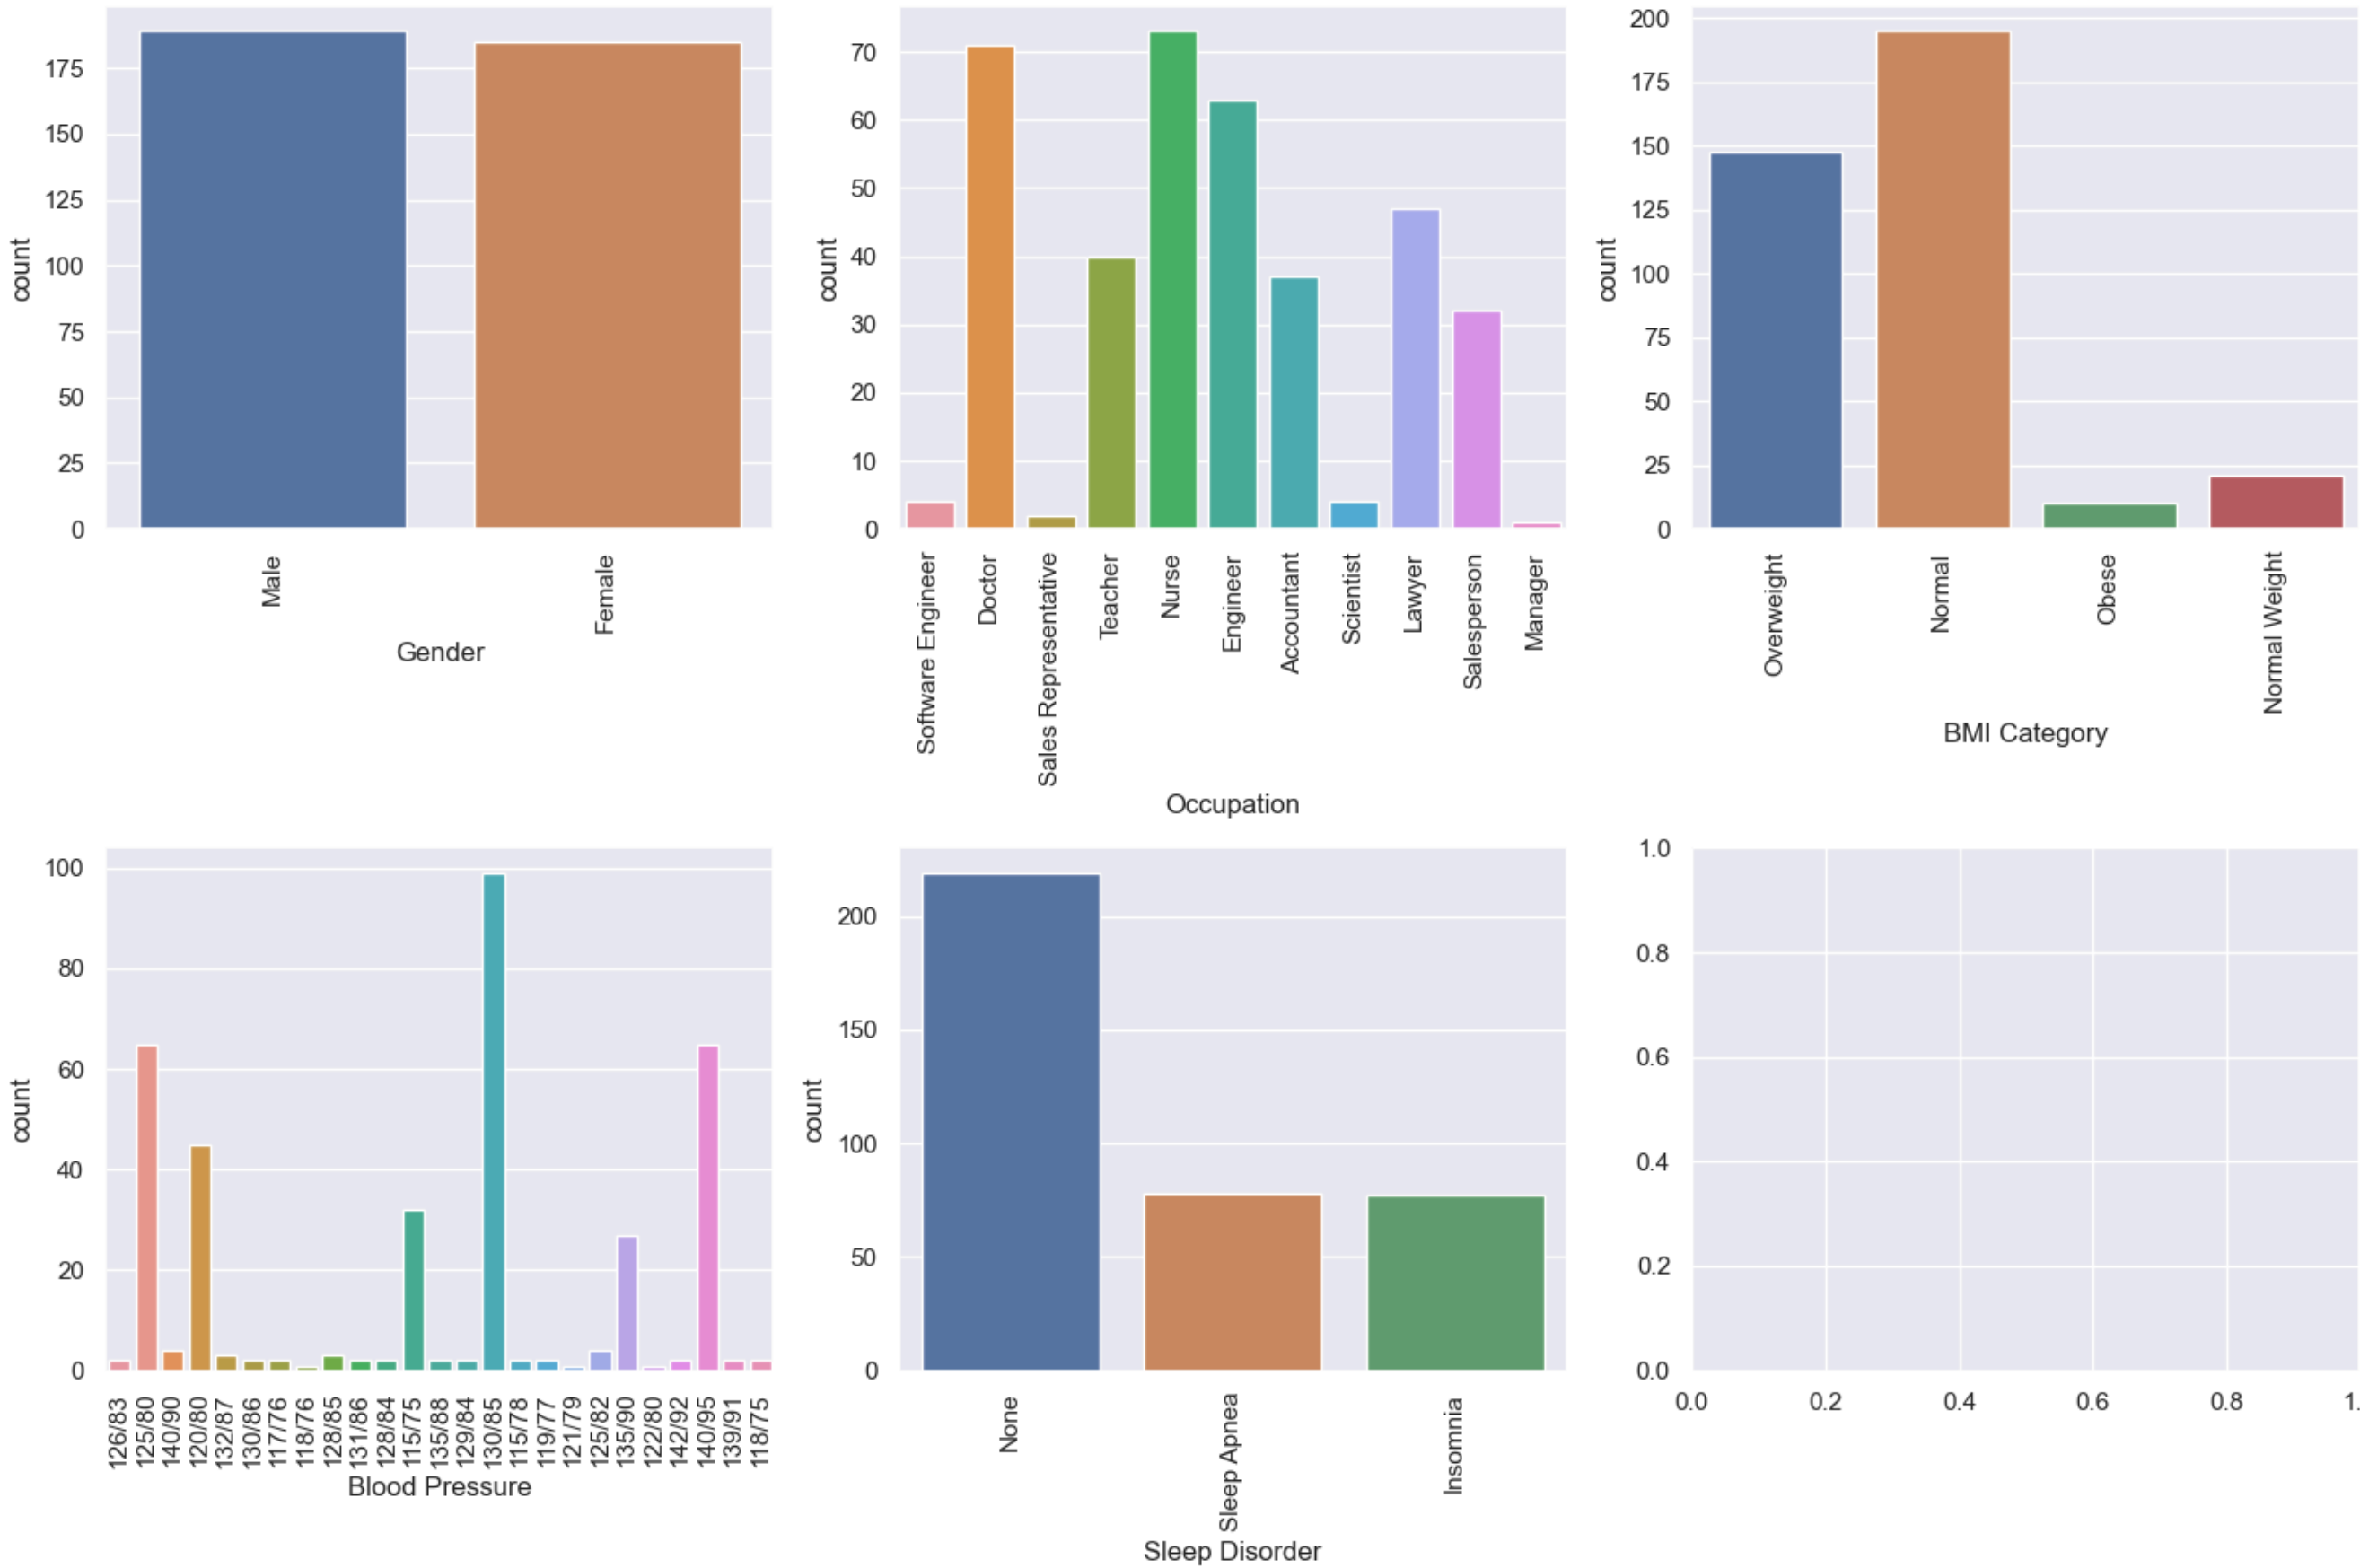
```
In [13]: # Listar as categorias para plotar
col_categ = ['Gender', 'Occupation', 'BMI Category', 'Blood Pressure', 'Sleep Disorder']

fig, ax = plt.subplots(nrows=2, ncols=3, figsize=(15,10))
axs = ax.flatten()


for i, var in enumerate(col_categ):
    sns.countplot(x=var, data=df_health, ax=axs[i])
    axs[i].set_xticklabels(axs[i].get_xticklabels(), rotation=90)

fig.tight_layout()

plt.show()
```

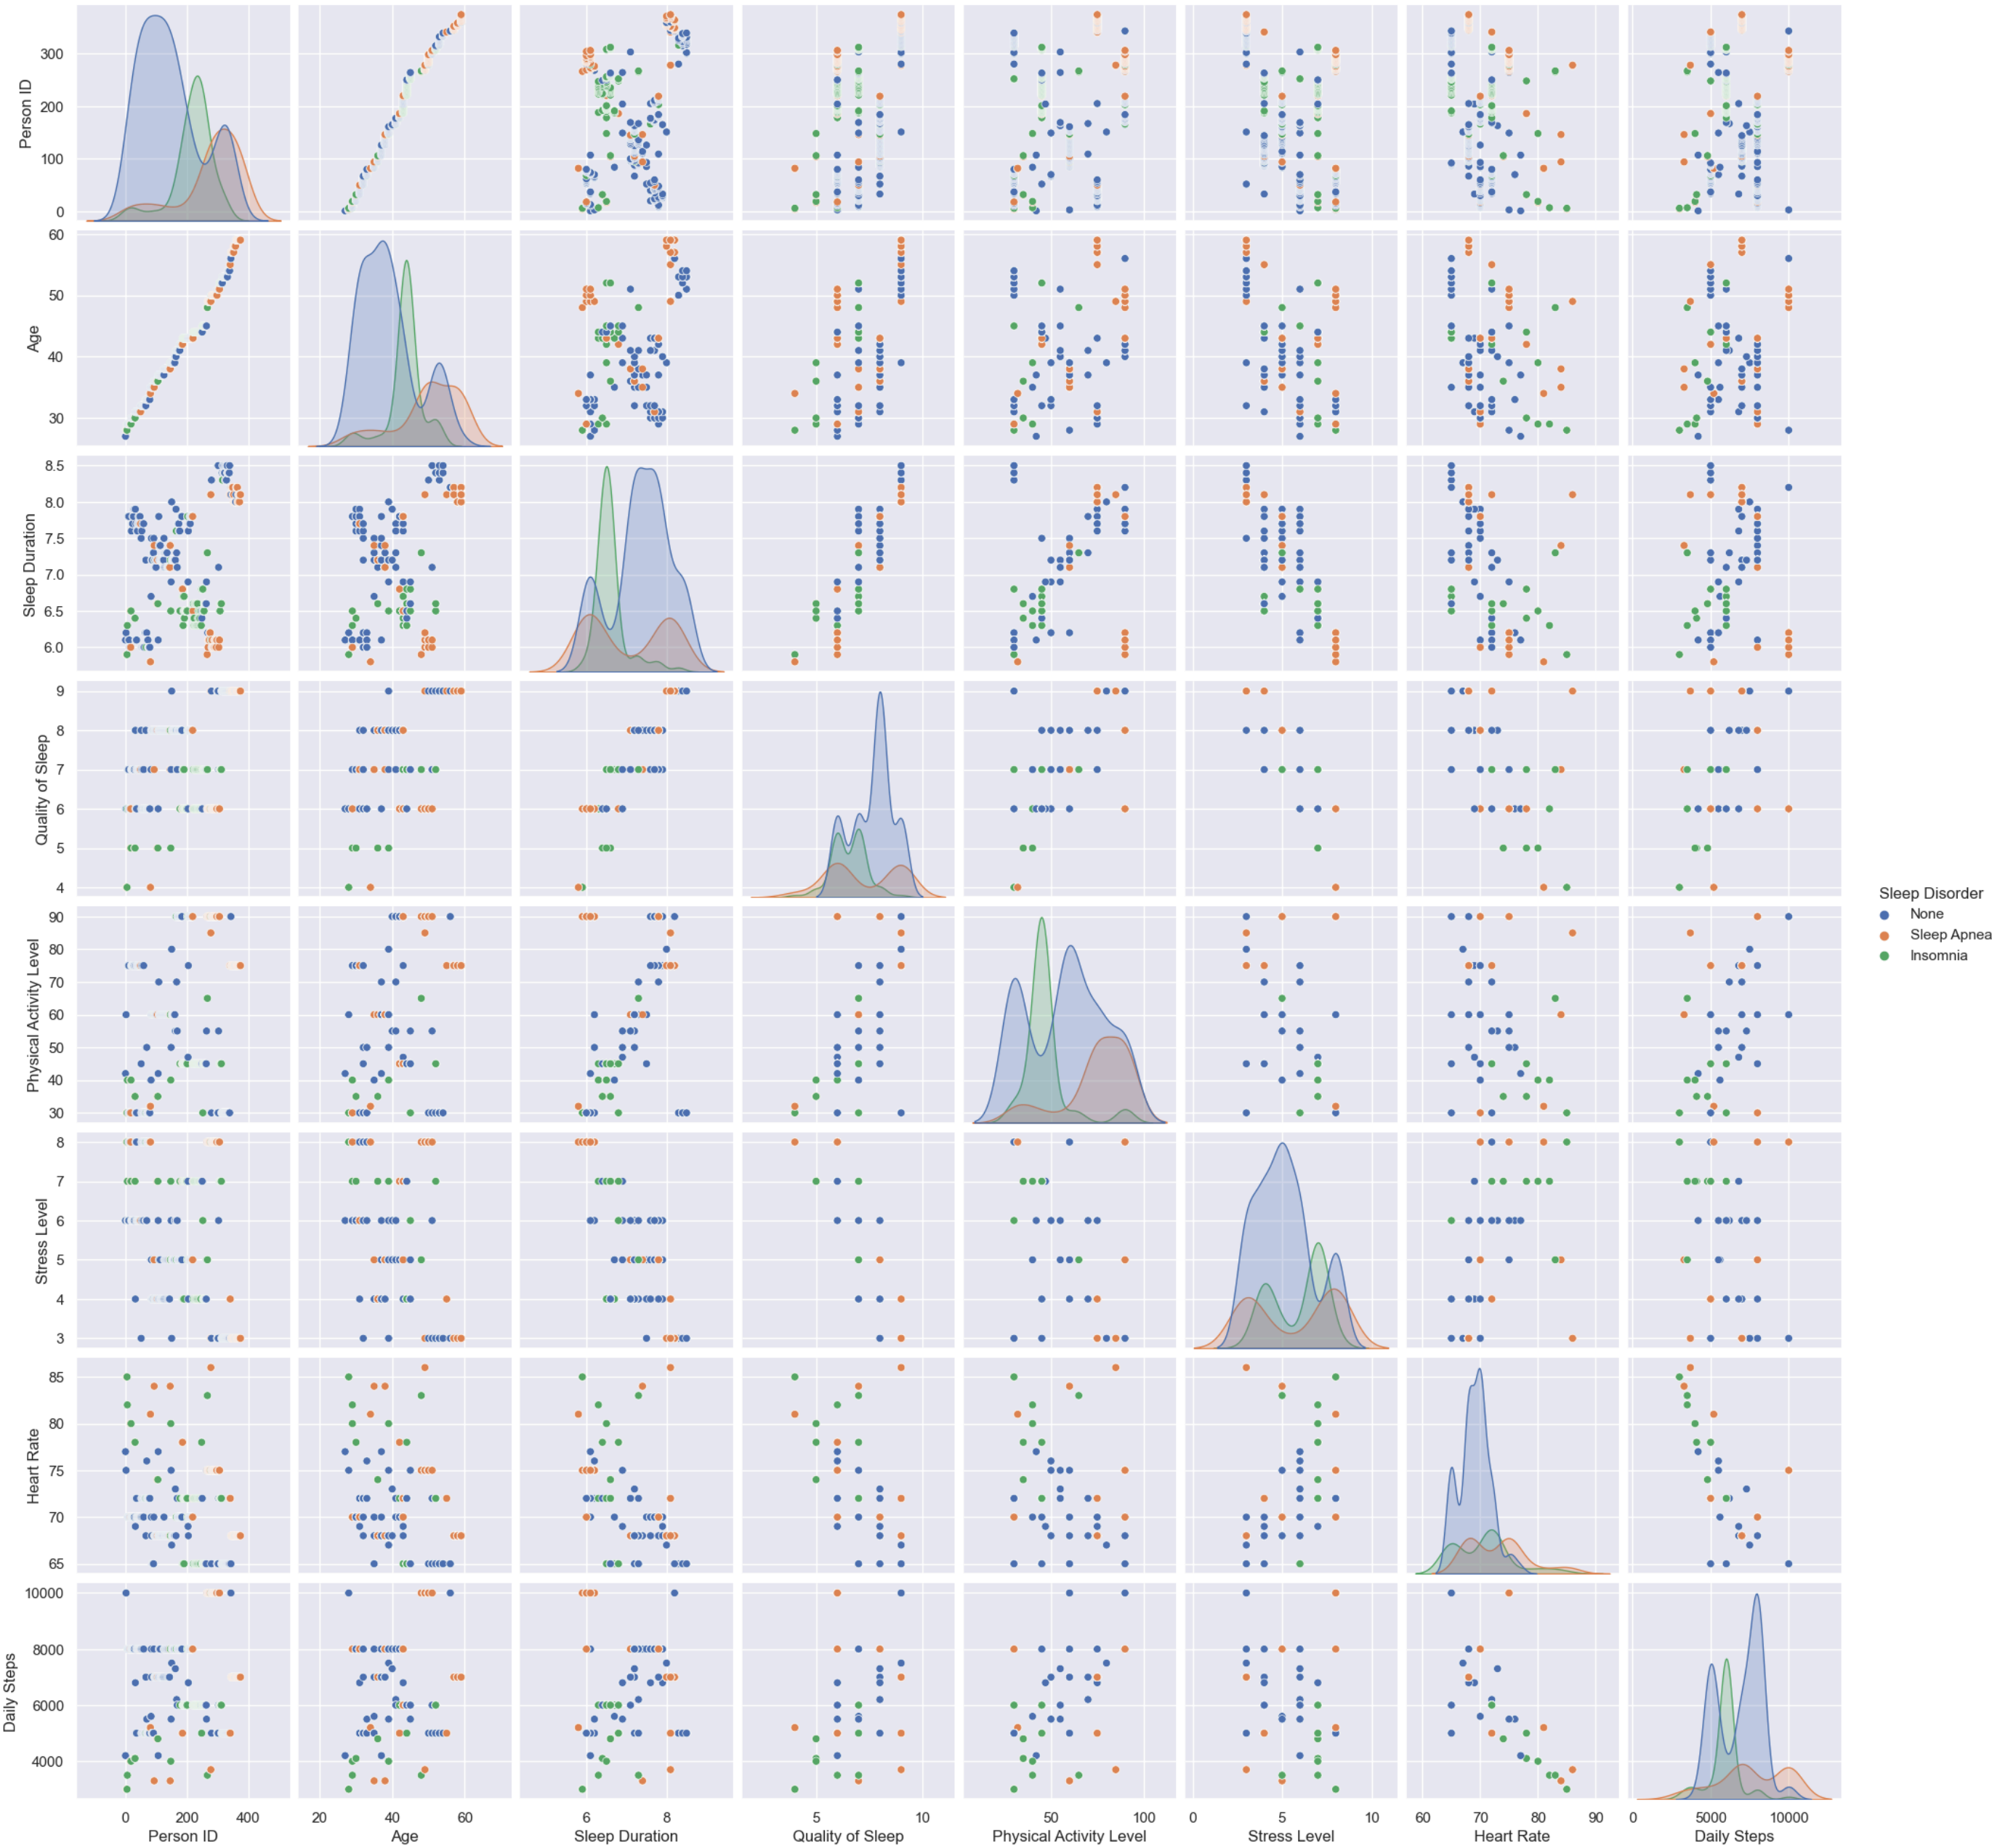


```
In [15]: sns.histplot(data=df_health, x='Age', kde=True);
```



```
In [16]: fig = plt.figure(figsize=(20,12))
sns.pairplot(data=df_health, hue='Sleep Disorder');
```

<Figure size 2000x1200 with 0 Axes>



Data Preprocessing Part 2

```
In [17]: df_health.head()
```

	Person ID	Gender	Age	Occupation	Sleep Duration	Quality of Sleep	Physical Activity Level	Stress Level	BMI Category	Blood Pressure	Heart Rate	Daily Steps	Sleep Disorder
0	1	Male	27	Software Engineer	6.1	6	42	6	Overweight	126/83	77	4200	None
1	2	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75	10000	None
2	3	Male	28	Doctor	6.2	6	60	8	Normal	125/80	75	10000	None
3	4	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85	3000	Sleep Apnea
4	5	Male	28	Sales Representative	5.9	4	30	8	Obese	140/90	85	3000	Sleep Apnea

```
In [18]: df_health['Blood Pressure'] = df_health['Blood Pressure'].str.replace('/', '.')
```

```
In [19]: ## Alterar as categorias para numeros
```

```
df_health['Gender'] = df_health['Gender'].apply(sexo)
df_health['BMI Category'] = df_health['BMI Category'].apply(BMI_Cat)
df_health['Sleep Disorder'] = df_health['Sleep Disorder'].apply(target)
```

```
In [20]: dummy_features = pd.get_dummies(df_health['Occupation'], drop_first=True)

In [21]: df_final = pd.concat([df_health, dummy_features], axis=1)

In [22]: df_final = df_final.drop('Occupation', axis=1)

In [23]: df_final.shape

Out[23]: (374, 22)
```

Split Train / Test

```
In [24]: X = df_final.drop('Sleep Disorder', axis=1)
         y = df_final['Sleep Disorder']

In [25]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

In [26]: X_train.shape, y_train.shape, X_test.shape, y_test.shape

Out[26]: ((261, 21), (261,), (113, 21), (113,))
```

Classifier Models

```
In [27]: # Logistic Regression
         model_lr = LogisticRegression()
         model_lr.fit(X_train, y_train)

Out[27]: ▾ LogisticRegression
         LogisticRegression()
```

```
In [28]: predict_lr = model_lr.predict(X_test)

In [29]: print(classification_report(y_test, predict_lr))

              precision    recall  f1-score   support

    0           0.66       0.82      0.73         62
    1           1.00       0.63      0.77         27
    2           0.37       0.29      0.33         24

 accuracy          0.66         113
 macro avg          0.68         113
weighted avg          0.68         113
```

```
In [30]: print(confusion_matrix(y_test, predict_lr))

[[51  0 11]
 [ 9 17  1]
 [17  0  7]]
```

```
In [31]: # Decision Tree Classifier
         model_dtc = DecisionTreeClassifier()
         model_dtc.fit(X_train, y_train)

Out[31]: ▾ DecisionTreeClassifier
         DecisionTreeClassifier()
```

```
In [32]: predict_dtc = model_dtc.predict(X_test)

In [33]: print(classification_report(y_test, predict_dtc))

              precision    recall  f1-score   support

    0           0.94       1.00      0.97         62
    1           0.95       0.74      0.83         27
    2           0.81       0.88      0.84         24

 accuracy          0.91         113
 macro avg          0.90         113
weighted avg          0.91         113
```

```
In [146... print(confusion_matrix(y_test, predict_dtc))

[[62  0  0]
 [ 2 20  5]
 [ 2  1 21]]
```

```
In [34]: # Random Forest Classifier
         model_rfc = RandomForestClassifier()
         model_rfc.fit(X_train, y_train)

Out[34]: ▾ RandomForestClassifier
         RandomForestClassifier()
```

```
In [35]: predict_rfc = model_rfc.predict(X_test)

In [36]: print(classification_report(y_test, predict_rfc))

              precision    recall  f1-score   support

    0           0.95       0.98      0.97         62
    1           0.91       0.74      0.82         27
    2           0.78       0.88      0.82         24

 accuracy          0.90         113
 macro avg          0.88         113
weighted avg          0.91         113
```

```
In [37]: print(confusion_matrix(y_test, predict_rfc))

[[62  0  0]
 [ 2 20  5]
 [ 2  1 21]]
```

```
In [38]: # Gradient Boosting Classifier
         model_hgb = HistGradientBoostingClassifier()
         model_hgb.fit(X_train, y_train)

Out[38]: ▾ HistGradientBoostingClassifier
         HistGradientBoostingClassifier()
```

```
In [39]: predict_hgb = model_hgb.predict(X_test)

In [40]: print(classification_report(y_test, predict_hgb))

              precision    recall  f1-score   support

    0           0.95       1.00      0.98         62
    1           0.91       0.74      0.82         27
    2           0.81       0.88      0.84         24

 accuracy          0.91         113
 macro avg          0.89         113
weighted avg          0.91         113
```

Comparação com a coluna Blood Pressure alterada

Logistic Regression

precision recall f1-score support				
0	0.75	0.81	0.78	62
1	0.95	0.74	0.83	27
2	0.48	0.50	0.49	24
accuracy			0.73	113
macro avg 0.73 0.68 0.70 113 weighted avg 0.74 0.73 0.73 113				

precision recall f1-score support				
0	0.66	0.82	0.73	62
1	1.00	0.63	0.77	27
2	0.37	0.29	0.33	24
accuracy			0.66	113
macro avg 0.68 0.58 0.61 113 weighted avg 0.68 0.66 0.66 113				

Decision Tree

precision recall f1-score support				
0	0.91	0.98	0.95	62
1	0.91	0.74	0.82	27
2	0.88	0.88	0.88	24
accuracy			0.90	113
macro avg 0.90 0.87 0.88 113 weighted avg 0.90 0.90 0.90 113				

precision recall f1-score support				
0	0.94	1.00	0.97	62
1	0.95	0.74	0.83	27
2	0.81	0.88	0.84	24
accuracy			0.91	113
macro avg 0.90 0.87 0.88 113 weighted avg 0.91 0.91 0.91 113				

Random Forest

precision recall f1-score support				
0	0.93	1.00	0.96	62
1	0.91	0.74	0.82	27
2	0.88	0.88	0.88	24
accuracy			0.91	113
macro avg 0.90 0.87 0.88 113 weighted avg 0.91 0.91 0.91 113				

precision recall f1-score support				
0	0.92	0.98	0.95	62
1	0.87	0.74	0.80	27
2	0.88	0.88	0.88	24
accuracy			0.90	113
macro avg 0.89 0.87 0.88 113 weighted avg 0.90 0.90 0.90 113				

Gradient Boosting Classifier

precision recall f1-score support				
0	0.95	1.00	0.98	62
1	0.91	0.74	0.82	27
2	0.81	0.88	0.84	24
accuracy			0.91	113
macro avg 0.89 0.87 0.88 113 weighted avg 0.91 0.91 0.91 113				